

Fall 2014

# Principles of Software Construction: Objects, Design, and Concurrency

## The Perils of Concurrency

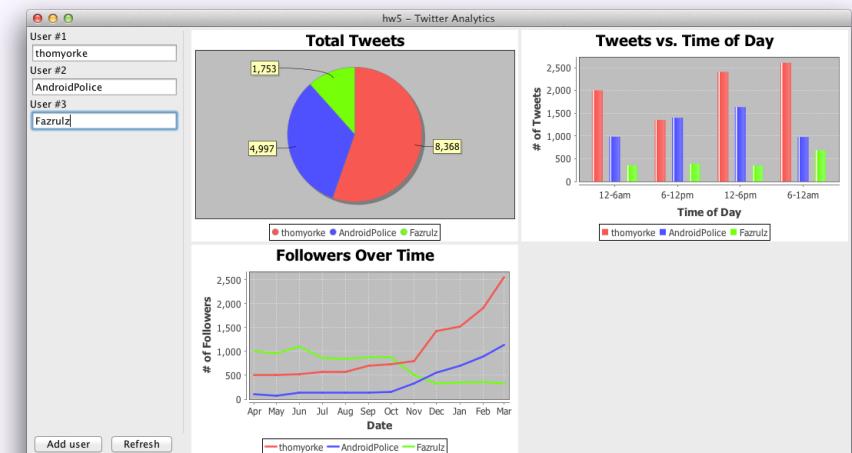
*Can't live with it.*

*Can't live without it.*

**Charlie Garrod   Jonathan Aldrich**

# Administrivia

- Homework 4c due tonight
  - Remember to add an ant run target
- 2<sup>nd</sup> midterm exam Thursday
  - Review session tonight, 5-7 p.m. in PH 100
- Homework 5 released by tomorrow
  - Must select partner(s) by Thursday (30 Oct)
  - 5a due next Wednesday morning (05 Nov)
  - 5b due the following Thursday (13 Nov)
  - 5c due the following Thursday (20 Nov)



# Key concepts from last Thursday

# API: Application Programming Interface

- An API defines the boundary between components/modules in a programmatic system

The screenshot shows a Java code editor with an XML configuration file. The code includes imports from org.omg.CORBA and com.ibm.ws.pmi.server, and various methods like size(), isEmpty(), iterator(), toArray(), and toArray(E[] a). The XML part defines a plugin with an extension for org.eclipse.ui.editors, pointing to a Sample XML Editor with XML as its extension and sample.gif as its icon. It also lists public repositories and provides parameters for type and sort.

```
org.omg.CORBA.MARSHAL: com.ibm.ws.pmi.server.DataDescriptor; IllegalAccessException minor code: 4942F23E comp
    at com.ibm.rmi.io.ValueHandlerImpl.readValue(ValueHandlerImpl.java:199)
    at com.ibm.rmi.iiop.CDRInputStream.read_value(CDRInputStream.java:1429)
    at com.ibm.rmi.io.ValueHandlerImpl.read_Array(ValueHandlerImpl.java:166)
    at com.ibm.rmi.io.ValueHandlerImpl.readValueInternal(ValueHandlerImpl.java:1429)
    at com.ibm.rmi.io.ValueHandlerImpl.readValue(ValueHandlerImpl.java:199)
    at com.ibm.rmi.iiop.CDRInputStream.read_value(CDRInputStream.java:1429)
    at com.ibm.ejs.sm.beans._EJSRemoteStatelessPmiService_Tie.<init>(_Tie.java:166)
    at com.ibm.CORBA.iiop.ExtendedServerDelegate.dispatch(ExtendedServerDelegate.java:186)
    at com.ibm.CORBA.iiop.ORB.process(ORB.java:2377)
    at com.ibm.CORBA.iiop.OrbWorker.run(OrbWorker.java:186)
    at com.ibm.ejs.oa.pool.ThreadPool$PooledWorker.run(ThreadPool.java:137)
    at com.ibm.ws.util.CachedThread.run(ThreadPool.java:137)

    int          size();
    boolean      isEmpty();
    Iterator<E> iterator();
    Object[]     toArray();
    E[]          toArray(E[] a);

AbstractLayoutCache.NodeDimensions
AbstractList
AbstractListModel
AbstractMap
AbstractMap.SimpleEntry
AbstractMap.SimpleImmutableEntry
AbstractMarshallerImpl
AbstractMethodError
AbstractOwnableSynchronizer

java.awt.im
java.awt.im.spi
java.awt.image
java.awt.image.renderable
java.awt.print

Provides classes and interfaces for ...
SortedMap<K, V>

<?xml version="1.0" encoding="UTF-8"?>
<?eclipse version="3.2"?>
<plugin>
<extension point="org.eclipse.ui.editors">
<editor name="Sample XML Editor" extensions="xml" icon="icons/sample.gif" contributorClass="org.eclipse.ui.texteditor.BasicTextEditorActionContributor" or="true" class="mveditor.editors.XMLEditor" id="mveditor.editors.XMLEditor">
</editor>
</extension>
</plugin>
```

# An API design process

- Define the scope of the API
  - Collect use-case stories, define requirements
  - Be skeptical
    - Distinguish true requirements from so-called solutions
    - "When in doubt, leave it out."
- Draft a specification, gather feedback, revise, and repeat
  - Keep it simple, short
- Code early, code often
  - Write *client code* before you implement the API

# Key design principle: Information hiding

- "When in doubt, leave it out."

# Minimize mutability

- Immutable objects are:
  - Inherently thread-safe
  - Freely shared without concern for side effects
  - Convenient building blocks for other objects
  - Can share internal implementation among instances
    - See `java.lang.String`
- Mutable objects require careful management of visibility and side effects
  - e.g. `Component.getSize()` returns a mutable `Dimension`
- Document mutability
  - Carefully describe state space



# The four course themes



- **Threads and concurrency**
  - Concurrency is a crucial system abstraction
  - E.g., background computing while responding to users
  - Concurrency is necessary for performance
  - Multicore processors and distributed computing
  - Our focus: application-level concurrency
  - Cf. functional parallelism (150, 210) and systems concurrency (213)
- **Object-oriented programming**
  - For flexible designs and reusable code
  - A primary paradigm in industry – basis for modern frameworks
  - Focus on Java – used in industry, some upper-division courses
- **Analysis and modeling**
  - Practical specification techniques and verification tools
  - Address challenges of threading, correct library usage, etc.
- **Design**
  - Proposing and evaluating alternatives
  - Modularity, information hiding, and planning for change
  - Patterns: well-known solutions to design problems

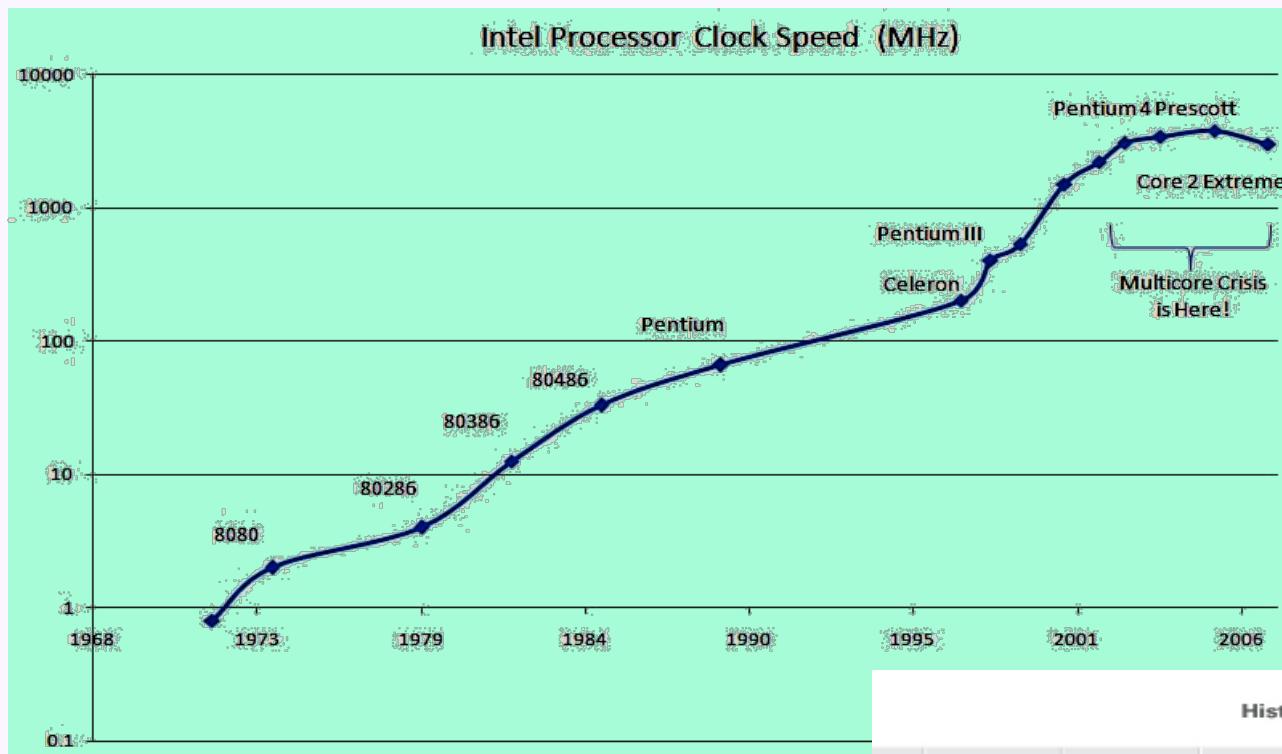
# Today: Concurrency, part 1

- The backstory
  - Motivation, goals, problems, ...
- Basic concurrency in Java
  - Synchronization
- Coming soon (but not today):
  - Higher-level abstractions for concurrency
    - Data structures
    - Computational frameworks

# Learning goals

- Understand concurrency as a source of complexity in software
- Know common abstractions for parallelism and concurrency, and the trade-offs among them
  - Explicit concurrency
    - Write thread-safe concurrent programs in Java
    - Recognize data race conditions
  - Know common thread-safe data structures, including high-level details of their implementation
  - Understand trade-offs between mutable and immutable data structures
  - Know common uses of concurrency in software design

# Processor speeds over time

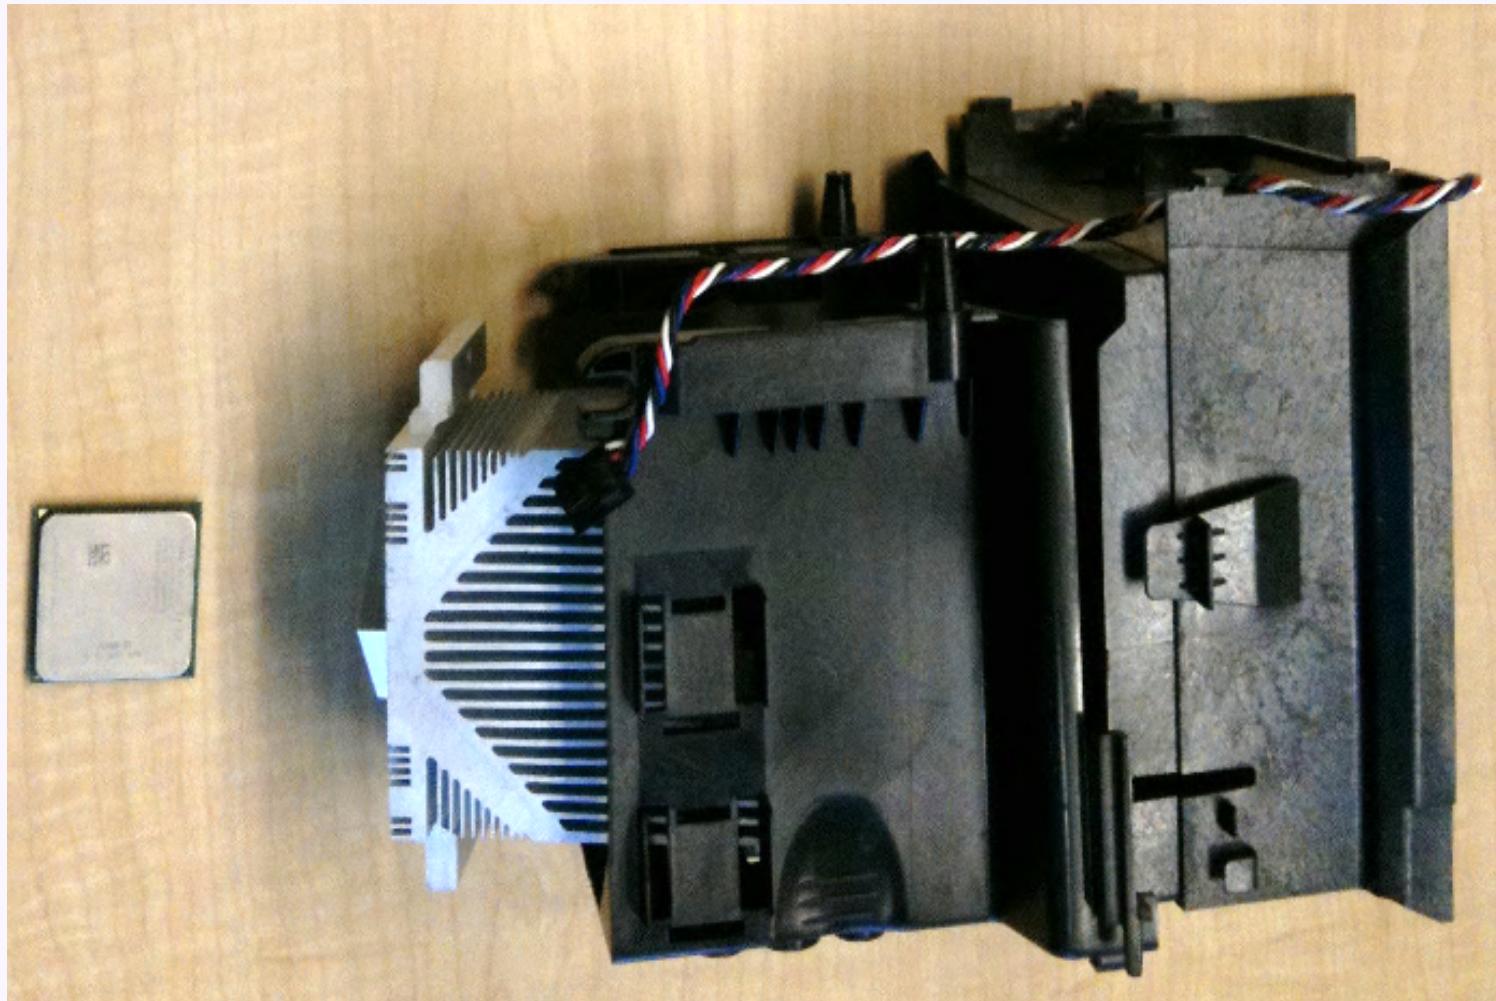


# Power requirements of a CPU

- Approx.: **Capacitance \* Voltage<sup>2</sup> \* Frequency**
- To increase performance:
  - More transistors, thinner wires: more **C**
    - More power leakage: increase **V**
  - Increase clock frequency **F**
    - Change electrical state faster: increase **V**
- Problem: Power requirements are super-linear to performance
  - Heat output is proportional to power input

# One option: fix the symptom

- Dissipate the heat



# One option: fix the symptom

- Better: Dissipate the heat with liquid nitrogen
  - Overclocking by Tom's Hardware's 5 GHz project



<http://www.tomshardware.com/reviews/5-ghz-project,731-8.html>

# Another option: fix the underlying problem

- Reduce heat by limiting power input
  - Adding processors increases power requirements linearly with performance
    - Reduce power requirement by reducing the frequency and voltage
    - Problem: requires concurrent processing

## Aside: Three sources of disruptive innovation

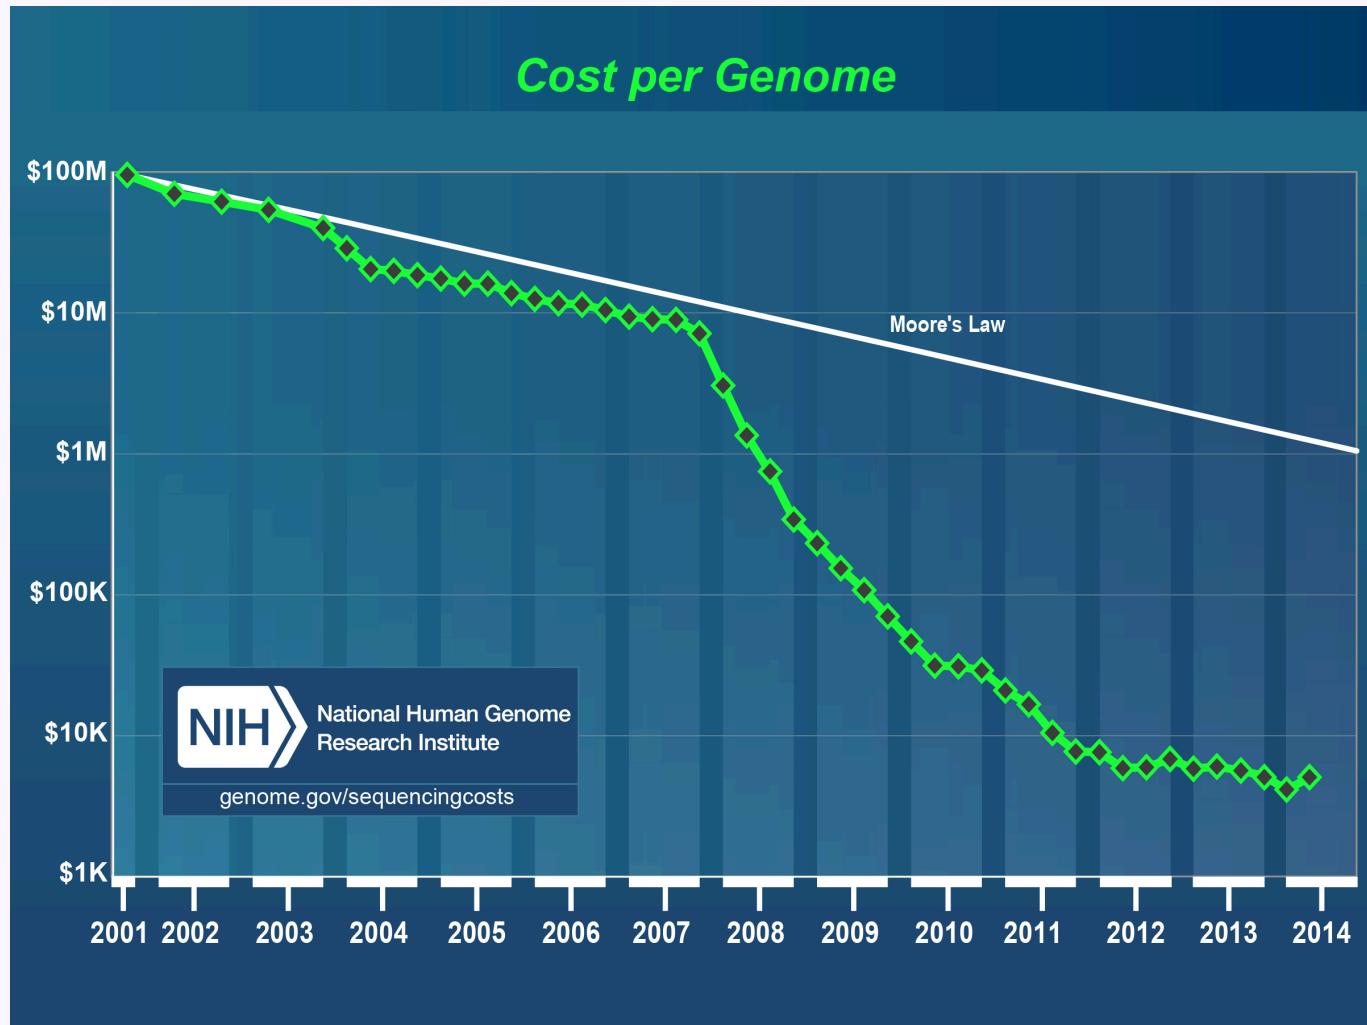
- Growth crosses some threshold
  - e.g., Concurrency: ability to add transistors exceeded ability to dissipate heat
- Colliding growth curves
  - Rapid design change forced by jump from one curve onto another
- Network effects
  - Amplification of small triggers leads to rapid change

## Aside: The threshold for distributed computing

- Too big for a single computer?
  - Forces use of distributed architecture
    - Shifts responsibility for reliability from hardware to software
      - Allows you to buy larger cluster of cheap flaky machines instead of expensive slightly-less-flaky machines
        - Revolutionizes data center design

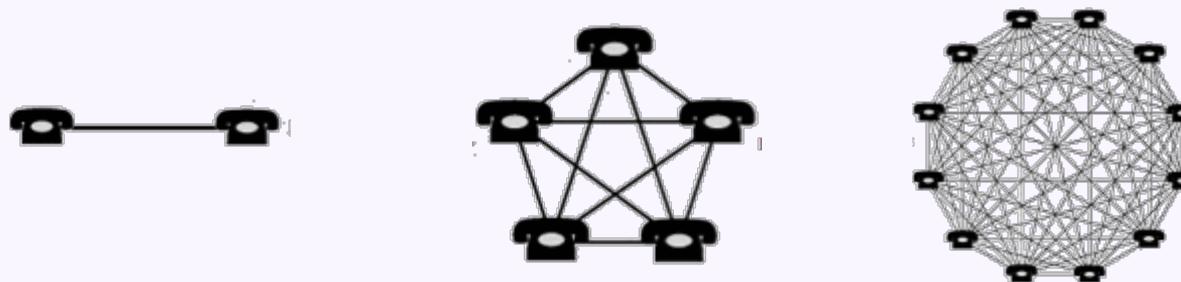
## Aside: Colliding growth curves

- From <http://www.genome.gov/sequencingcosts/>



## Aside: Network effects

- Metcalfe's rule: network value grows quadratically in the number of nodes
  - a.k.a. Why my mom has a Facebook account
  - $n(n-1)/2$  potential connections for  $n$  nodes



- Creates a strong imperative to merge networks
  - Communication standards, USB, media formats, ...

# Concurrency

- Simply: doing more than one thing at a time
  - In software: more than one point of control
    - Threads, processes
- Resources simultaneously accessed by more than one thread or process

# Concurrency then and now

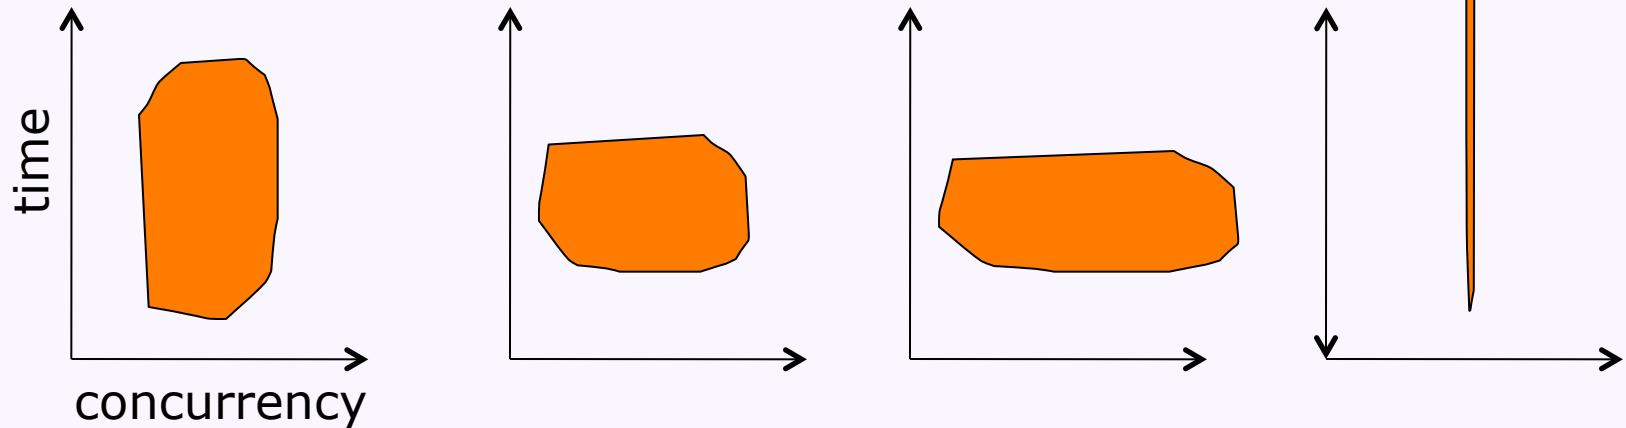
- In the past multi-threading was just a convenient abstraction
  - GUI design: event threads
  - Server design: isolate each client's work
  - Workflow design: producers and consumers
- Now: must use concurrency for scalability and performance

Image Name	Threads	C
IPSSVC.EXE	86	0
svchost.exe	82	0
System	80	0
afsd_service.exe	51	0
Rtvscan.exe	47	0
winlogon.exe	39	0
explorer.exe	20	0
ccEvtMgr.exe	19	0
svchost.exe	18	0
lsass.exe	18	0
tabtip.exe	17	0
svchost.exe	17	0
firefox.exe	16	0
services.exe	16	0
thunderbird.exe	15	0
csrss.exe	13	0
tcserver.exe	10	0
KeyboardSurroga...	10	0
spoolsv.exe	10	0
tvt_reg_monitor_...	10	0
svchost.exe	10	0
POWERPNT.EXE	9	0
taskmgr.exe	8	0
VPTray.exe	8	0
S24EvMon.exe	8	0
EvtEng.exe	8	0
emacs.exe	7	0
tvtsched.exe	7	0
ibmpmsvc.exe	7	0
AcroRd32.exe	7	0
vpngui.exe	6	0
cvpnd.exe	6	0
AluSchedulerSvc....	6	0
ccSetMgr.exe	6	0
svchost.exe	6	0
wisptis.exe	5	0
alg.exe	5	0
TPHKMGR.exe	5	0
ASRSVC.exe	5	0

# Problems of concurrency

- Realizing the potential
  - Keeping all threads busy doing useful work
- Delivering the right language abstractions
  - How do programmers think about concurrency?
  - Aside: parallelism vs. concurrency
- Non-determinism
  - Repeating the same input can yield different results

# Realizing the potential



- Possible metrics of success
  - Breadth: extent of simultaneous activity
    - width of the shape
  - Depth (or span): length of longest computation
    - height of the shape
  - Work: total effort required
    - area of the shape
- Typical goals in parallel algorithm design?

# Amdahl's law: How good can the depth get?

- Ideal parallelism with  $N$  processors:

- Speedup =  $N$

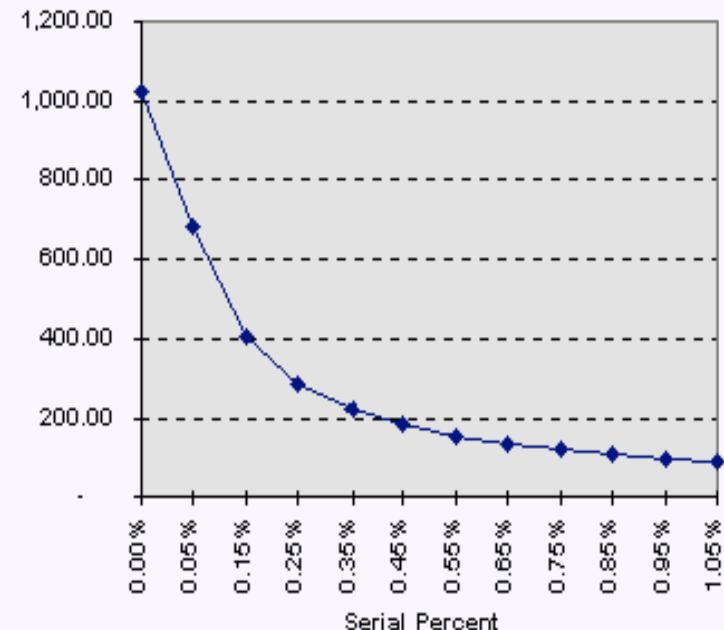
- In reality, some work is always inherently sequential

- Let  $F$  be the portion of the total task time that is inherently sequential

- Speedup = 
$$\frac{1}{F + (1 - F)/N}$$

- Suppose  $F = 10\%$ . What is the max speedup? (you choose  $N$ )

Speedup by Amdahl's Law ( $P=1024$ )



# Amdahl's law: How good can the depth get?

- Ideal parallelism with  $N$  processors:

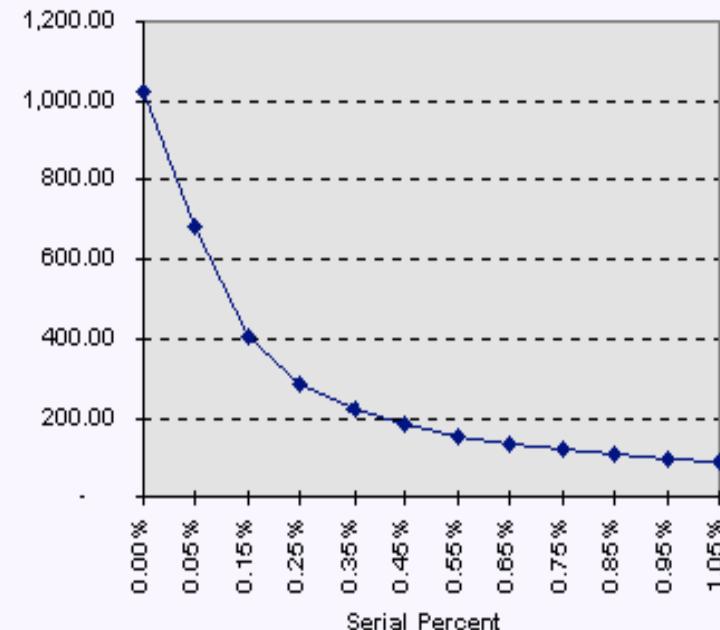
- Speedup =  $N$

- In reality, some work is always inherently sequential

- Let  $F$  be the portion of the total task time that is inherently sequential

- Speedup = 
$$\frac{1}{F + (1 - F)/N}$$

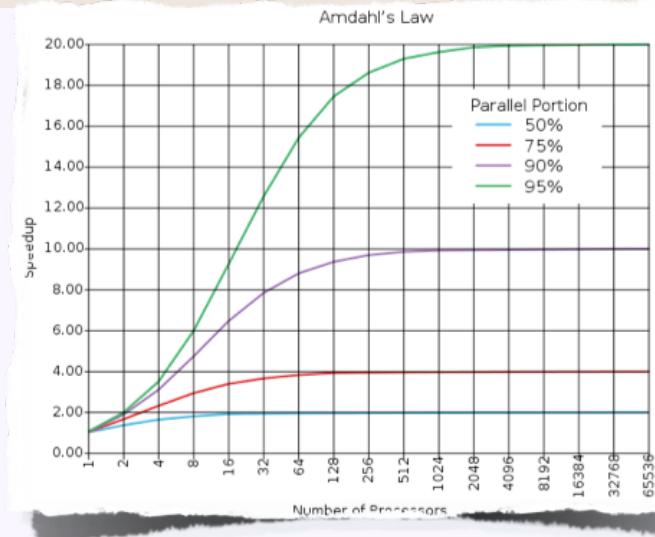
Speedup by Amdahl's Law ( $P=1024$ )



- Suppose  $F = 10\%$ . What is the max speedup? (you choose  $N$ )
    - As  $N$  approaches  $\infty$ ,  $1/(0.1 + 0.9/N)$  approaches 10.

# Using Amdahl's law as a design guide

- For a given algorithm, suppose
  - $N$  processors
  - Problem size  $M$
  - Sequential portion  $F$
- An obvious question:
  - What happens to speedup as  $N$  scales?
- A less obvious, important question:
  - What happens to  $F$  as problem size  $M$  scales?

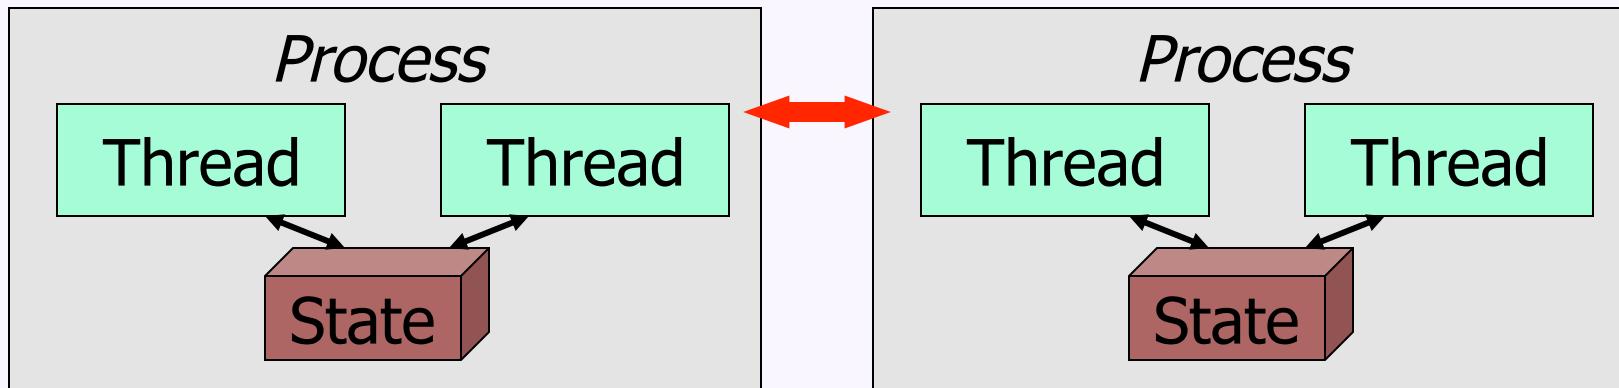


*"For the past 30 years, computer performance has been driven by Moore's Law; from now on, it will be driven by Amdahl's Law."*

— Doron Rajwan, Intel Corp

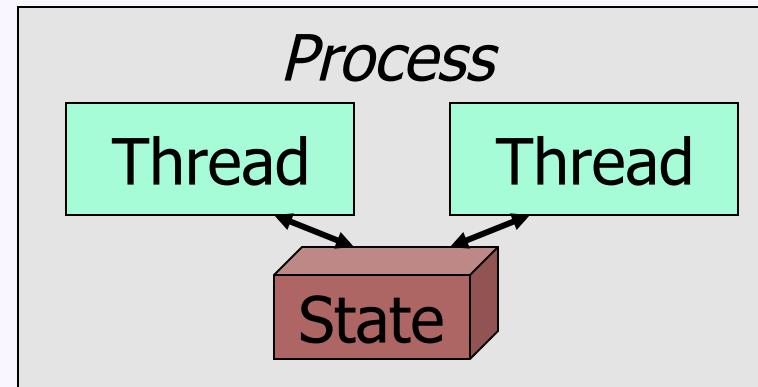
# Abstractions of concurrency

- **Processes**
  - Execution environment is isolated
    - Processor, in-memory state, files, ...
  - Inter-process communication typically slow, via message passing
    - Sockets, pipes, ...
- **Threads**
  - Execution environment is shared
  - Inter-thread communication typically fast, via shared state

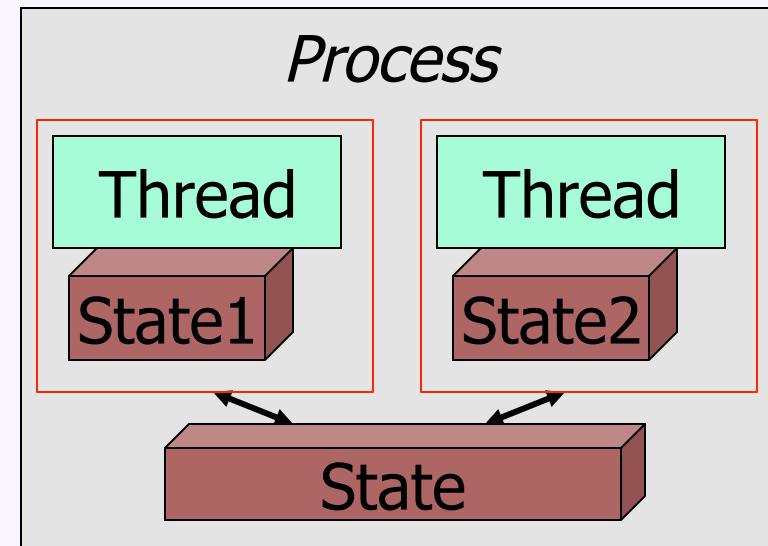


## Aside: Abstractions of concurrency

- What you see:
  - State is all shared



- A (slightly) more accurate view of the hardware:
  - Separate state stored in registers and caches
  - Shared state stored in caches and memory



# Basic concurrency in Java

- The `java.lang.Runnable` interface

```
void run();
```

- The `java.lang.Thread` class

```
Thread(Runnable r);
void start();
static void sleep(long millis);
void join();
boolean isAlive();
static Thread currentThread();
```

- See `IncrementTest.java`

# Atomicity

- An action is *atomic* if it is indivisible
  - Effectively, it happens all at once
    - No effects of the action are visible until it is complete
    - No other actions have an effect during the action
- In Java, integer increment is not atomic

i++;

is actually

1. Load data from variable i
2. Increment data by 1
3. Store data to variable i

# One concurrency problem: race conditions

- A *race condition* is when multiple threads access shared data and unexpected results occur depending on the order of their actions
- E.g., from `IncrementTest.java`:
  - Suppose `classData` starts with the value 41:

Thread A:

```
classData++;
```

Thread B:

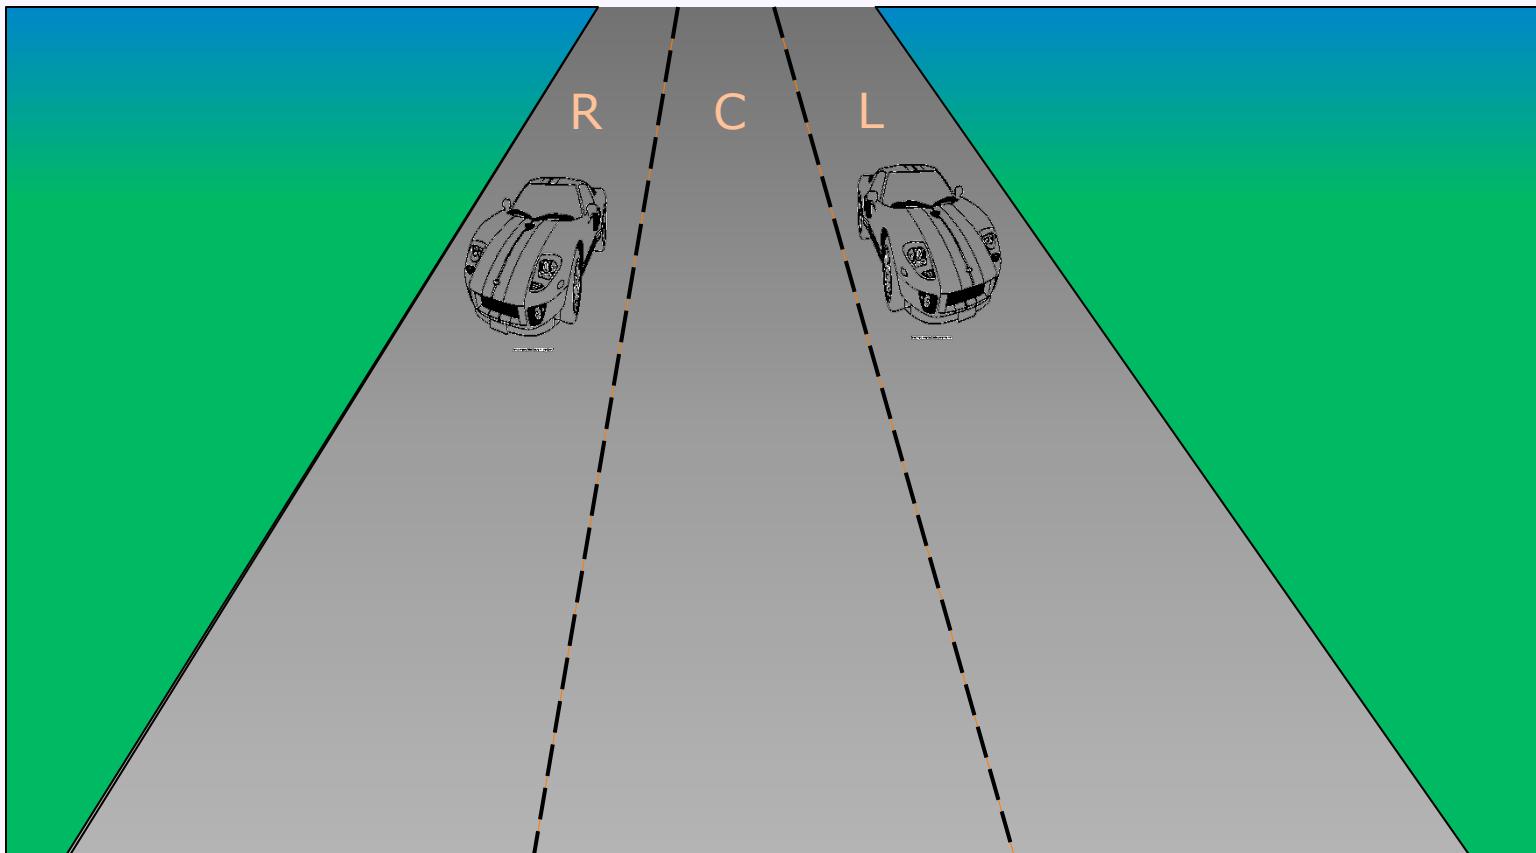
```
classData++;
```

One possible interleaving of actions:

- 1A. Load data(41) from `classData`
- 1B. Load data(41) from `classData`
- 2A. Increment data(41) by 1 -> 42
- 2B. Increment data(41) by 1 -> 42
- 3A. Store data(42) to `classData`
- 3B. Store data(42) to `classData`

# Race conditions in real life

- E.g., check-then-act on the highway



# Race conditions in real life

- E.g., check-then-act at the bank
  - The "debit-credit problem"

## *Alice, Bob, Bill, and the Bank*

- A. *Alice* to pay *Bob* \$30
  - Bank actions
    1. Does Alice have \$30 ?
    2. Give \$30 to *Bob*
    3. Take \$30 from *Alice*
- B. *Alice* to pay *Bill* \$30
  - Bank actions
    1. Does Alice have \$30 ?
    2. Give \$30 to *Bill*
    3. Take \$30 from *Alice*
- If *Alice* starts with \$40, can *Bob* and *Bill* both get \$30?

# Race conditions in real life

- E.g., check-then-act at the bank
  - The "debit-credit problem"

## *Alice, Bob, Bill, and the Bank*

- A. *Alice* to pay *Bob* \$30
  - Bank actions
    1. Does Alice have \$30 ?
    2. Give \$30 to *Bob*
    3. Take \$30 from *Alice*
- B. *Alice* to pay *Bill* \$30
  - Bank actions
    1. Does Alice have \$30 ?
    2. Give \$30 to *Bill*
    3. Take \$30 from *Alice*
- If *Alice* starts with \$40, can *Bob* and *Bill* both get \$30?

A.1  
A.2  
B.1  
B.2  
A.3  
B.3!

# Race conditions in *your* real life

- E.g., check-then-act in simple code

```
public class StringConverter {  
    private Object o;  
    public void set(Object o) {  
        this.o = o;  
    }  
    public String get() {  
        if (o == null) return "null";  
        return o.toString();  
    }  
}
```

- See StringConverter.java, Getter.java, Setter.java

# Some actions are atomic

Precondition:

```
int i = 7;
```

Thread A:

```
i = 42;
```

Thread B:

```
ans = i;
```

- What are the possible values for ans?

# Some actions are atomic

Precondition:

```
int i = 7;
```

Thread A:

```
i = 42;
```

Thread B:

```
ans = i;
```

- What are the possible values for ans?

```
i: 00000...00000111
```

:

```
i: 00000...00101010
```

# Some actions are atomic

Precondition:

```
int i = 7;
```

Thread A:

```
i = 42;
```

Thread B:

```
ans = i;
```

- What are the possible values for ans?

i: **00000...00000111**

:

i: **00000...00101010**

- In Java:

- Reading an int variable is atomic
- Writing an int variable is atomic

- Thankfully, ans: **00000...00101111**

is not possible

# Bad news: some simple actions are not atomic

- Consider a single 64-bit long value

**high bits**

**low bits**

- Concurrently:
  - Thread A writing high bits and low bits
  - Thread B reading high bits and low bits

Precondition:

```
long i = 1000000000;
```

Thread A:

```
i = 42;
```

Thread B:

```
ans = i;
```

ans: **01001...00000000**

(1000000000)

ans: **00000...00101010**

(42)

ans: **01001...00101010**

(1000000042 or ...)

# Thursday:

- Midterm exam...
- Next week:
  - Primitive concurrency control
  - ...then higher abstractions