

Principles of Software Construction: Objects, Design, and Concurrency

Distributed System Design, Part 2

Fall 2014

Charlie Garrod Jonathan Aldrich

Administrivia

- Homework 5b due tonight
 - Finish by tomorrow (14 Nov) 10 a.m. if you want to be considered as a "Best Framework" for Homework 5c
- 15-413: Software Engineering Practicum
- Homework 3 arena winners in class next week...

Key concepts from Tuesday

Networking in Java

- The `java.net.InetAddress`:

```
static InetAddress getByName(String host);
static InetAddress getByAddress(byte[ ] b);
static InetAddress getLocalHost();
```


- The `java.net.Socket`:

```
Socket(InetAddress addr, int port);
boolean isConnected();
boolean isClosed();
void close();
InputStream getInputStream();
OutputStream getOutputStream();
```

- The `java.net.ServerSocket`:

```
ServerSocket(int port);
Socket accept();
void close();
...
```

Aside: The robustness vs. redundancy curve



Metrics of success

- Reliability
 - Often in terms of availability: fraction of time system is working
 - 99.999% available is "5 nines of availability"
- Scalability
 - Ability to handle workload growth

Today: Distributed system design

- Introduction to distributed systems, continued
 - Motivation: reliability and scalability
 - Failure models
 - Techniques for:
 - Reliability (availability)
 - Scalability
 - Consistency
- MapReduce: A robust, scalable framework for distributed computation...
 - ...on replicated, partitioned data

Types of failure behaviors

- Fail-stop
- Other halting failures
- Communication failures
 - Send/receive omissions
 - Network partitions
 - Message corruption
- Data corruption
- Performance failures
 - High packet loss rate
 - Low throughput
 - High latency
- Byzantine failures

Common assumptions about failures


- Behavior of others is fail-stop (ugh)
- Network is reliable (ugh)
- Network is semi-reliable but asynchronous
- Network is lossy but messages are not corrupt
- Network failures are transitive
- Failures are independent
- Local data is not corrupt
- Failures are reliably detectable
- Failures are unreliably detectable

Some distributed system design goals


- The end-to-end principle
 - When possible, implement functionality at the ends (rather than the middle) of a distributed system
- The robustness principle
 - Be strict in what you send, but be liberal in what you accept from others
 - Protocols
 - Failure behaviors
- Benefit from incremental changes
- Be redundant
 - Data replication
 - Checks for correctness

Replication for scalability: Client-side caching

- Architecture before replication:




- Problem: Server throughput is too low
- Solution: Cache responses at (or near) the client
 - Cache can respond to repeated read requests




Replication for scalability: Client-side caching

- Hierarchical client-side caches:




Replication for scalability: Server-side caching

- Architecture before replication:



- Problem: Database server throughput is too low
- Solution: Cache responses on multiple servers
 - Cache can respond to repeated read requests



Cache invalidation


- Time-based invalidation (a.k.a. expiration)
 - Read-any, write-one
 - Old cache entries automatically discarded
 - No expiration date needed for read-only data
- Update-based invalidation
 - Read-any, write-all
 - DB server broadcasts invalidation message to all caches when the DB is updated

Cache replacement policies

- Problem: caches have finite size
- Common* replacement policies
 - Optimal (Belady's) policy
 - Discard item not needed for longest time in future
 - Least Recently Used (LRU)
 - Track time of previous access, discard item accessed least recently
 - Least Frequently Used (LFU)
 - Count # times item is accessed, discard item accessed least frequently
 - Random
 - Discard a random item from the cache

Partitioning for scalability

- Partition data based on some property, put each partition on a different server




Horizontal partitioning

- a.k.a. "sharding"
- A table of data:

username	school	value
cohen	CMU	9
bob	CMU	42
alice	Yale	90
pete	Yale	12
deb	MIT	16
reif	MIT	40


Recall: Basic hash tables

- For n-size hash table, put each item x in the bucket: $x.\text{hashCode}() \% n$



Partitioning with a distributed hash table

- Each server stores data for one bucket
- To store or retrieve an item, front-end server hashes the key, contacts the server storing that bucket



Consistent hashing


- Goal: Benefit from incremental changes
 - Resizing the hash table (i.e., adding or removing a server) should not require moving many objects
- E.g., Interpret the range of hash codes as a ring
 - Each bucket stores data for a range of the ring
 - Assign each bucket an ID in the range of hash codes
 - To store item x don't compute $x.hashCode() \% n$. Instead, place x in bucket with the same ID as or next higher ID than $x.hashCode()$

Problems with hash-based partitioning

- Front-ends need to determine server for each bucket
 - Each front-end stores look-up table?
 - Master server storing look-up table?
 - Routing-based approaches?
- Places related content on different servers
 - Consider *range* queries:
`SELECT * FROM users WHERE lastname STARTSWITH 'G'`

Master/tablet-based systems

- Dynamically allocate range-based partitions
 - Master server maintains tablet-to-server assignments
 - Tablet servers store actual data
 - Front-ends cache tablet-to-server assignments




Today: Distributed system design

- Introduction to distributed systems, continued
 - Motivation: reliability and scalability
 - Failure models
 - Techniques for:
 - Reliability (availability)
 - Scalability
 - Consistency
- MapReduce: A robust, scalable framework for distributed computation...
 - ...on replicated, partitioned data

Map from a functional perspective

- `map(f, x[0..n-1])`
 - Apply the function f to each element of list x



map/reduce images src: Apache Hadoop tutorials

- E.g., in Python:

```
def square(x): return x*x  
map(square, [1, 2, 3, 4]) would return [1, 4, 9, 16]
```

- Parallel map implementation is trivial
 - What is the work? What is the depth?

Reduce from a functional perspective


- `reduce(f, x[0...n-1])`

- Repeatedly apply binary function f to pairs of items in x , replacing the pair of items with the result until only one item remains
- One sequential Python implementation:

```
def reduce(f, x):  
    if len(x) == 1: return x[0]  
    return reduce(f, [f(x[0],x[1])] + x[2:])
```


- e.g., in Python:

```
def add(x,y): return x+y  
reduce(add, [1,2,3,4])  
    would return 10 as  
reduce(add, [1,2,3,4])  
reduce(add, [3,3,4])  
reduce(add, [6,4])  
reduce(add, [10]) -> 10
```



Reduce with an associative binary function

- If the function f is associative, the order f is applied does not affect the result



$$1 + ((2+3) + 4) \quad 1 + (2 + (3+4)) \quad (1+2) + (3+4)$$

- Parallel reduce implementation is also easy
 - What is the work? What is the depth?

Distributed MapReduce


- The distributed MapReduce idea is similar to (but not the same as!):

```
reduce(f2, map(f1, x))
```

- Key idea: a "data-centric" architecture
 - Send function f_1 directly to the data
 - Execute it concurrently
 - Then merge results with reduce
 - Also concurrently
- Programmer can focus on the data processing rather than the challenges of distributed systems

MapReduce with key/value pairs (Google style)

- Master
 - Assign tasks to workers
 - Ping workers to test for failures
- Map workers
 - Map for each key/value pair
 - Emit intermediate key/value pairs
- Reduce workers
 - Sort data by intermediate key and aggregate by key
 - Reduce for each key



MapReduce with key/value pairs (Google style)

- E.g., for each word on the Web, count the number of times that word occurs
 - For Map: key1 is a document name, value is the contents of that document
 - For Reduce: key2 is a word, values is a list of the number of counts of that word

```
f1(String key1, String value):  
    for each word w in value:  
        EmitIntermediate(w, 1);
```

```
f2(String key2, Iterator values):  
    int result = 0;  
    for each v in values:  
        result += v;  
    Emit(key2, result);
```

Map: $(key1, v1) \rightarrow (key2, v2)^*$


Reduce: $(key2, v2^*) \rightarrow (key3, v3)^*$

MapReduce: $(key1, v1)^* \rightarrow (key3, v3)^*$

MapReduce: $(docName, docText)^* \rightarrow (word, wordCount)^*$

MapReduce architectural details


- Usually integrated with a distributed storage system
 - Map worker executes function on its share of the data
- Map output usually written to worker's local disk
 - Shuffle: reduce worker often pulls intermediate data from map worker's local disk
- Reduce output usually written back to distributed storage system



Handling server failures with MapReduce

- Map worker failure:
 - Re-map using replica of the storage system data
- Reduce worker failure:
 - New reduce worker can pull intermediate data from map worker's local disk, re-reduce
- Master failure:

- Options:
 - Restart system using new master
 - Replicate master
 - ...



The beauty of MapReduce

- Low communication costs (usually)
 - The shuffle (between map and reduce) is expensive
- MapReduce can be iterated
 - Input to MapReduce: key/value pairs in the distributed storage system
 - Output from MapReduce: key/value pairs in the distributed storage system

Another MapReduce example

- E.g., for person in a social network graph, output the number of mutual friends they have
 - For Map: key1 is a person, value is the list of her friends
 - For Reduce: key2 is ???, values is a list of ???

```
f1(String key1, String value):
```

```
f2(String key2, Iterator values):
```

MapReduce: (person, friends)* → (pair of people, count of mutual friends)*

Another MapReduce example

- E.g., for person in a social network graph, output the number of mutual friends they have
 - For Map: key1 is a person, value is the list of her friends
 - For Reduce: key2 is a pair of people, values is a list of 1s, for each mutual friend that pair has

```
f1(String key1, String value):  
    for each pair of friends  
        in value:  
            EmitIntermediate(pair, 1);
```

```
f2(String key2, Iterator values):  
    int result = 0;  
    for each v in values:  
        result += v;  
    Emit(key2, result);
```

MapReduce: (person, friends)* → (pair of people, count of mutual friends)*

And another MapReduce example

- E.g., for each page on the Web, create a list of the pages that link to it
 - For Map: key1 is a document name, value is the contents of that document
 - For Reduce: key2 is ???, values is a list of ???

f1(String key1, String value):

f2(String key2, Iterator values):

MapReduce: (docName, docText)* \rightarrow (docName, list of incoming links)*

Coming next...

- More distributed systems
 - MapReduce